

I am an **Interactive Machine Translation** researcher at Lilt with expertise in **Data Science**, **Machine Learning**, **Natural Language Processing**, and **Human-Computer Interaction**. I have a Ph.D in Computer Science from Stanford University, and have been awarded the NSF and NDSEG Fellowships, and a Stanford Human-Centered AI Grant. I have written software used by millions of users – I built HabitLab, a data-driven behavior change system with **12,000+ daily active users**; I also wrote and maintain UNetbootin, an open-source project with **40 million downloads**.

Education

Ph.D. , Computer Science	Stanford University	GPA: 4.0/4.0	2019
B.S. and M.Eng. , Computer Science	Massachusetts Institute of Technology	GPA: 5.0/5.0	2013

Work Experience

Principal Research Scientist	Lilt	San Francisco	Feb 2021 - present
Senior Research Scientist	Lilt	San Francisco	Aug 2019 - Feb 2021

I am the head of HCI research and manage a team of researchers at Lilt, an interactive machine translation startup. Improved interactive MT system speed by shifting computation client-side via TensorflowJS and heuristics. Built named entity transliteration system based on Transformer architecture using tensor2tensor and Tensorflow. Developed metrics and logging to determine how translators spend their time and predict translator performance. Ran A/B tests to evaluate website translation ROI, and developed system that recommends pages to translate.

Research Intern	Microsoft Research	Redmond	Summer 2015
------------------------	--------------------	---------	-------------

Designed and built an educational social feed app usable by pre-literate children. Published at CSCW 2017.

Research Intern	Microsoft Research	Beijing	Summer 2014
------------------------	--------------------	---------	-------------

Built QuizCram, a quiz-driven MOOC lecture viewer that improves learning outcomes. Presented at CHI 2015.

Software Engineering Intern	Google	Mountain View	Summer 2013
------------------------------------	--------	---------------	-------------

Developed a machine learning system for detecting taps on the phone bezel, for use in Android input methods.

Software Engineering Intern	Google	Mountain View	Summer 2012
------------------------------------	--------	---------------	-------------

Developed an NLP system to automatically generate glossaries from book text. Patent granted **US9483460B2**.

Software Engineering Intern	Google	Mountain View	Summer 2011
------------------------------------	--------	---------------	-------------

Developed a machine learning system to predict the quality of user reviews, now deployed on Google Play.

Software Development Engineer Intern	Microsoft	Redmond	Summer 2010
---	-----------	---------	-------------

Built the IntelliSense code completion system for a scientific computing language, and contributed to its compiler.

Skills and Technologies

Programming Languages: Python, JavaScript, C, C++, Java, TypeScript, R, C#, Ruby, Scala, Haskell, Bash, SQL
Machine Learning + Deep Learning: Tensorflow, PyTorch, TensorflowJS, Keras, scikit-learn, xgboost, MLFlow
Natural Language Processing + Machine Translation: SpaCy, tensor2tensor, fairseq, HuggingFace Transformers
Data Science + Visualization: NumPy, SciPy, Pandas, Jupyter, RStudio, Plotly, Superset, Spark, Hadoop MapReduce
Web + Mobile Development: HTML, CSS, React, AngularJS, NodeJS, Express, D3.js, Flask, MySQL, Docker, Android
Languages: Fluent: English, Chinese (Mandarin), Hungarian. Intermediate: Japanese, Vietnamese, Spanish.

Publications in Journals and Conference Proceedings

Geza Kovacs, Zhengxuan Wu, Michael Bernstein. "Not Now, Ask Later: Users Weaken Their Behavior Change Regimen Over Time, But Expect To Re-Strengthen It Imminently." *ACM annual conference on Human Factors in Computing Systems (CHI) 2021*. Acceptance rate: 23%.

Samuel Läubli, Patrick Simianer, Joern Wuebker, **Geza Kovacs**, Rico Sennrich, Spence Green. "The Impact of Text Presentation on Translator Performance." *Target: International Journal of Translation Studies, 2021 (to appear)*.

Geza Kovacs, Drew Mylander Gregory, Zilin Ma, Zhengxuan Wu, Golrokh Emami, Jacob Ray, Michael Bernstein. "Conservation of Procrastination: Do Productivity Interventions Save Time Or Just Redistribute It?" *ACM annual conference on Human Factors in Computing Systems (CHI) 2019*. Acceptance rate: 23.8%.

Geza Kovacs, Zhengxuan Wu, Michael Bernstein. "Rotating Online Behavior Change Interventions Increases Effectiveness But Also Increases Attrition." *ACM annual conference on Computer-Supported Cooperative Work and Social Computing (CSCW) 2018*. Acceptance rate: 26%.

Rajan Vaish, Neil Gaiwad, **Geza Kovacs**, Andreas Veit, Ranjay Krishna, Imanol Arrieta Ibarra, Camelia Simoiu, Michael Wilber, Serge Belongie, Sharad Goel, James Davis, Michael Bernstein. "Crowd Research: Open and Scalable University Laboratories." *ACM Symposium on User Interface Software and Technology (UIST) 2017*. Acceptance rate: 22%.

Kiley Sobel, **Geza Kovacs**, Galen McQuillen, Andrew Cross, Nirupama Chandrasekaran, Nathalie Riche, Ed Cutrell, Meredith Morris. "EduFeed: A Social Feed to Engage Preliterate Children in Educational Activities." *ACM annual conference on Computer-Supported Cooperative Work and Social Computing (CSCW) 2017*. Acceptance rate: 35%.

Geza Kovacs. "Effects of In-Video Quizzes on MOOC Lecture Viewing." *ACM annual conference on Learning at Scale (L@S) 2016*. Acceptance rate: 22%.

Geza Kovacs and Robert C. Miller. "Smart Subtitles for Vocabulary Learning." *ACM annual conference on Human Factors in Computing Systems (CHI) 2014*. Acceptance rate: 23%.

Publications in Adjunct Conference Proceedings

Stanford Crowd Research, **Geza Kovacs**, Rajan Vaish, Michael Bernstein. "Daemo: A Self-Governed Crowdsourcing Marketplace." *ACM Symposium on User Interface Software and Technology (UIST) 2015, Poster*.

Geza Kovacs. "FeedLearn: Using Facebook Feeds for Microlearning." *ACM annual conference on Human Factors in Computing Systems (CHI) 2015, Extended Abstracts*.

Geza Kovacs. "QuizCram: A Question-Driven Video Studying Interface." *ACM annual conference on Human Factors in Computing Systems (CHI) 2015, Extended Abstracts*.

Joseph Jay Williams, **Geza Kovacs**, Caren Walker, Samuel G Maldonado, Tania Lombrozo. "Learning Online via Prompts to Explain." *ACM annual conference on Human Factors in Computing Systems (CHI) 2014, Extended Abstracts*.

Geza Kovacs and Robert C. Miller. "Foreign Manga Reader: Learn Grammar and Pronunciation while Reading Comics." *ACM Symposium on User Interface Software and Technology (UIST) 2013, Demo*.

Geza Kovacs. "Smart Subtitles for Language Learning." *ACM annual conference on Human Factors in Computing Systems (CHI) 2013, Extended Abstracts*.

Geza Kovacs. "ScreenMatch: providing context to software translators by displaying screenshots." *ACM annual conference on Human Factors in Computing Systems (CHI) 2012, Extended Abstracts*.

Patents

Tania Bedrax-Weiss, **Geza Kovacs**, Ulas Kirazci. "Automated formation of specialized dictionaries." US9483460B2. Filed 10/2013, Published 11/2016, Expires 01/2034.

Meredith Morris, Nathalie Henry Riche, Edward B. Cutrell, Andrew C. Cross, Natasa Milic, Nirupama Chandrasekaran, Galen McQuillen, Kiley Sobel, **Geza Kovacs**. "Presenting educational activities via an extended social media feed." US20180068578A1. Filed 09/2016, Published 03/2018.

Open Source Projects

UNetbootin (LiveUSB Creator) <https://en.wikipedia.org/wiki/UNetbootin>
40 million downloads. UNetbootin creates bootable USB flash drives for various (50+) Linux distributions.

Wubi (Ubuntu Installer for Windows) [https://en.wikipedia.org/wiki/Wubi_\(software\)](https://en.wikipedia.org/wiki/Wubi_(software))
Now part of Ubuntu. Built the first versions of Wubi, which allows Ubuntu to be installed from Windows.

HabitLab (In-the-wild Behavior Change Research Platform) <https://habitlab.stanford.edu>
12,000+ daily active users. I built HabitLab over my Ph.D, and it is still used for research at Stanford Medical School.

Invited Keynote Talks

Geza Kovacs. "Predictive Translation Memory in the Wild: A Study of Interactive Machine Translation Use on Lilt." *Association for Machine Translation in the Americas (AMTA) Workshop on the Impact of Machine Translation 2020.*

Select Awards and Honors

Stanford Human-Centered AI Grant (for my research project HabitLab)	2018
National Defense Science and Engineering Graduate Fellowship	2013
National Science Foundation Graduate Research Fellowship	2013
1 st place, ACM UIST (User Interface Software and Technology) Student Innovation Contest	2012
1 st place, ACM CHI (Human Factors in Computing Systems) Student Research Competition	2012
Phi Beta Kappa (top 10% of students at MIT), Tau Beta Pi (top 12.5% of Engineering students at MIT)	2012

Academic Conference Reviewing and Committees

Organizing Committee, WMT 2022 Shared Task on Word-Level Auto-Completion	2021
Program Committee, EACL 2021 Bridging HCI and NLP Workshop	2021
Reviewer, ACM Conference on Human Factors in Computing Systems (CHI)	2015, 2018, 2019, 2021, 2022
Reviewer, ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW)	2021
Reviewer, ACM Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)	2019
Reviewer, ACM Symposium on User Interface Software and Technology (UIST)	2017, 2018

Researchers Managed

Sai Gouravajhala, Senior Research Scientist at Lilt.	August 2020 – present
Hannah Yan, Senior Data Scientist at Lilt.	September 2020 – present
Jordan Huffaker, Research Intern at Lilt. Now a Ph.D student at University of Michigan.	Summer 2021
Jessy Lin, Research Engineer at Lilt. Now a Ph.D student at UC Berkeley.	August 2019 – August 2020
Ming-Chang Chiu, Data Science Intern at Lilt. Now a Ph.D student at USC.	Summer 2020

Teaching Experience

Understanding Users (CS 377U) – Teaching Assistant, at Stanford	Spring 2019
Human Computer Interaction Research (CS 376) – Teaching Assistant, at Stanford	Fall 2018
Natural Language Processing (6.863) – Teaching Assistant, at MIT	Fall 2012

Select Coursework

Deep Learning (Stanford CS230), Natural Language Processing (MIT 6.864+6.863), Data Science (Stanford CS224w), Machine Learning (MIT 6.034), Statistical Models (MIT 6.804), Statistics (MIT 18.440), Linear Algebra (MIT 18.700), UX Design (MIT 6.803+MAS.672), Linguistics (MIT 24.900), Bioinformatics (MIT 6.047), Algorithms (MIT 6.006+6.046)

Select Press Coverage

HabitLab

WIRED - The HabitLab Browser Extension Curbs Your Time Wasted on the Web.	January 2019
Lifehacker - Prevent Procrastination With This Chrome Extension.	February 2019
The New York Times - Finding It Hard to Focus? Maybe It's Not Your Fault.	August 2018
Lifehacker - Be More Mindful of the Time You Waste Online With HabitLab.	August 2018
Entrepreneur - Use These Strategies to Maximize Productivity Without Inventing an Extra Weekday.	May 2018

Crowd Research / Daemo

Stanford University News - A Stanford-led platform for crowdsourced research	October 2017
WIRED - Amazon's Turker Crowd Has Had Enough.	August 2017

UNetbootin

Forbes - How To Try Linux Without Making Any Changes To Your PC.	September 2018
PCWorld - Create a Bootable Linux Flash Drive in Three Easy Steps.	February 2012
Lifehacker - The Complete Guide to Saving Your Windows System with a Thumb Drive.	March 2010
Network World - Installing Ubuntu on an old netbook with hair tearing and profanity.	August 2014

Wubi

Ars Technica - Wubi arrives: a look at Ubuntu 8.04 alpha 5.	February 2008
Lifehacker - Install Ubuntu on a Windows Netbook, No Partitioning Needed.	May 2008