Geza Kovacs

geza@cs.stanford.edu gkovacs.com

Education

Stanford University	PhD	Computer Science	GPA: 4.0/4.0	*2013 – now*
Massachusetts Institute of Technology	BS+MEng	Computer Science	GPA: 5.0/5.0	*2008 – 2013*

Industry Experience

Microsoft Research – Research Intern, Redmond – published at CSCW 2017 (EduFeed) *Summer 2015*
Microsoft Research – Research Intern, Beijing – published at CHI EA 2015 (QuizCram) *Summer 2014*
Google – Software Engineering Intern, Mountain View *Summer 2013*
Developed a machine learning system for detecting taps on the phone bezel, for use in Android input methods.
Google – Software Engineering Intern, Mountain View *Summer 2012*
Developed an NLP model to detect vocabulary and generate glossaries from book text (used MapReduce).
Google – Software Engineering Intern, Mountain View *Summer 2011*
Developed a machine learning model to predict the quality of user reviews of Android apps.
Microsoft Corporation, Redmond – Software Development Engineer Intern *Summer 2010*
Google Summer of Code – FFmpeg (open-source video transcoding library) *Summer 2009*

Research Highlights

HabitLab: Large-scale Online Behavior Change Experiments – published at CHI 2019 and CSCW 2018
HabitLab is an online experimentation platform with *12,000+ daily active users* that I developed during my PhD at Stanford. I have used it to conduct a variety of experiments, data science, and machine learning work:

- Predicted changes in users' intervention preferences over time (using **LSTM networks**; Python/PyTorch)
- Analyzed time redistribution effects caused by interventions (using **mixed models**; R/Python/SciPy)
- Analyzed effects of rotating interventions on effectiveness and attrition (**cox regression and LMM**; R)
- Personalized interventions to each user based on effectiveness (using **reinforcement learning**; Python)
- Predicted time spent on webpages, based on browsing visit history data (using **random forests**; Python/H2O)

Effects of In-Video Quizzes on MOOC Lecture Viewing – published at Learning at Scale 2016

- A large-scale data analysis of Coursera's in-video interaction logs across Machine Learning courses, analyzing effects of in-video quizzes on users' video watching and seeking behaviors (Python/Pandas)

Open-source Projects

UNetbootin (LiveUSB Creator) – http://unetbootin.github.io/ https://en.wikipedia.org/wiki/UNetbootin
40 million downloads. UNetbootin creates bootable USB flash drives for various (50+) Linux distributions.

Ubuntu Installer for Windows (Wubi) https://en.wikipedia.org/wiki/Wubi_(software)
Now part of Ubuntu. Built the first versions of Wubi, which allows Ubuntu to be installed from Windows.

Teaching Experience

Understanding Users (CS 377U) at Stanford – Teaching Assistant *Spring 2019*
Human Computer Interaction Research (CS 376) at Stanford – Teaching Assistant *Fall 2018*
Natural Language Processing (6.863) at MIT – Teaching Assistant *Fall 2012*

Relevant Coursework

Deep Learning (CS 230), **Natural Language Processing** (6.864+6.863), AI (6.034), Network Analysis (CS 224W), Computational Cognitive Science (6.804), Computational Biology (6.047), HCI (6.803), Computer Security (6.857), Compilers (CS 143), Algorithms (6.006+6.046), Linear Algebra (18.700), Probability (18.440)

Skills and technologies

Machine Learning: PyTorch, sklearn, Keras, TensorFlow, H2O, RL, Deep Learning (RNN/LSTM/CNN/GAN)
Data Mining: NumPy, SciPy, Pandas, NLTK, NetworkX, MapReduce, CUDA, SQL, NoSQL (MongoDB/Redis)
Data Science: Mixed models, Survival analysis, Experiment design, A/B testing, Multi-armed bandits, NLP
Data Visualization: D3.js, ggplot2, Plotly, Bokeh, Chartjs, matplotlib, Jupyter, RStudio
Web Development: HTML/CSS/JS, Node.js, Flask, Polymer, React, Flow, MongoDB, PostgreSQL, EC2
Mobile Development: Cross-platform JS (Cordova, NativeScript), Android (Java), Responsive Web Design
Programming Languages: Python, JavaScript, R, Java, C, C++, C#, Scala, Ruby, CoffeeScript, Haskell, Bash

Awards and Honors

Stanford Human-Centered AI Grant (for my research project HabitLab), 2018
National Defense Science and Engineering Graduate Fellowship, 2013
National Science Foundation Graduate Research Fellowship, 2013
1st place, Most Useful, ACM UIST (User Interface Software and Technology) Student Innovation Contest, 2012
1st place, ACM CHI (Conference on Human Factors in Computing Systems) Student Research Competition, 2012
1st place, MIT Autonomous Robotics Competition (Maslab), 2010

Journal and Conference Papers

Geza Kovacs, Drew Mylander Gregory, Zilin Ma, Zhengxuan Wu, Golrokh Emami, Jacob Ray, Michael Bernstein. "Conservation of Procrastination: Do Productivity Interventions Save Time Or Just Redistribute It?" ACM annual conference on Human Factors in Computing Systems (CHI) 2019.

Geza Kovacs, Zhengxuan Wu, Michael Bernstein. "Rotating Online Behavior Change Interventions Increases Effectiveness But Also Increases Attrition." ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW) 2018.

Rajan Vaish, Neil Gaikwad, **Geza Kovacs**, Andreas Veit, Ranjay Krishna, Imanol Arrieta Ibarra, Camelia Simoiu, Michael Wilber, Serge Belongie, Sharad Goel, James Davis, Michael Bernstein. "Crowd Research: Open and Scalable University Laboratories." ACM Symposium on User Interface Software and Technology (UIST) 2017.

Kiley Sobel, **Geza Kovacs**, Galen McQuillen, Andrew Cross, Nirupama Chandrasekaran, Nathalie Riche, Ed Cutrell, Meredith Morris. "EduFeed: A Social Feed to Engage Preliterate Children in Educational Activities." ACM annual conference on Computer Supported Collaborative Work (CSCW) 2017.

Geza Kovacs. "Effects of In-Video Quizzes on MOOC Lecture Viewing." ACM annual conference on Learning at Scale (L@S) 2016.

Geza Kovacs and Robert C. Miller. "Smart Subtitles for Vocabulary Learning." ACM annual conference on Human Factors in Computing Systems (CHI) 2014.

Peer-Reviewed Extended Abstracts

Stanford Crowd Research, **Geza Kovacs**, Rajan Vaish, Michael Bernstein. "Daemo: A Self-Governed Crowdsourcing Marketplace". ACM Symposium on User Interface Software and Technology (UIST) 2015, Poster.

Geza Kovacs. "FeedLearn: Using Facebook Feeds for Microlearning." ACM annual conference on Human Factors in Computing Systems (CHI) 2015, Extended Abstracts.

Geza Kovacs. "QuizCram: A Question-Driven Video Studying Interface." ACM annual conference on Human Factors in Computing Systems (CHI) 2015, Extended Abstracts.

Joseph Jay Williams, **Geza Kovacs**, Caren Walker, Samuel G Maldonado, Tania Lombrozo. "Learning Online via Prompts to Explain." ACM annual conference on Human Factors in Computing Systems (CHI) 2014, Extended Abstracts.

Geza Kovacs and Robert C. Miller. "Foreign Manga Reader: Learn Grammar and Pronunciation while Reading Comics." ACM Symposium on User Interface Software and Technology (UIST) 2013, Demo.

Geza Kovacs. "Smart Subtitles for Language Learning." ACM annual conference on Human Factors in Computing Systems (CHI) 2013, Extended Abstracts.

Geza Kovacs. "ScreenMatch: providing context to software translators by displaying screenshots." ACM annual conference on Human Factors in Computing Systems (CHI) 2012, Extended Abstracts.